

TensorFlow Implementation of Yolo9000

Tristan Partridge
V00804280

1

resulting shape of the output is $13 \times 13 \times 125$ and can be seen in Figure 2.

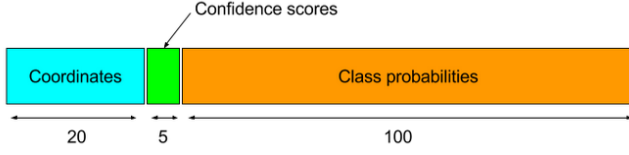


Figure 2. The output of the Yolo Network [2].

2.2. Anchor Boxes

Anchor boxes are an innovative aspect of how the Yolo9000 network makes detection predictions. One possible way a network could make bounding box predictions is by predicting the center x , center y , width, and height of each object. However, this approach is prone to overfitting and will take a very long time to train. The width and height of objects are very inconsistent and not conducive to consistent, accurate predictions for images the network has never seen before. Yolo gets around this by using *anchor boxes*. Each cell in the 13×13 grid has 5 anchors associated with it. These anchors are centered on the cell and are defined by a width and height relative to the size of the cell, each anchor having a different aspect ratio. For each image the network needs to learn the x and y offset from the center of the ground truth box to the center of the cell, as well as the width and height relative to the anchor box width and height. The exponential of the width and height ratios and sigmoid of x and y offsets are actually used as they are easier for the network to learn. Figure 3 shows an example of a ground truth bounding box offset from an anchor. When training, the anchor box with the lowest intersection over union (IOU) is associated with each ground truth bounding box. This technique of using anchor boxes keeps the bounding box predictions that the network needs to make much more consistent with each other and is more robust to a large variation in object sizes.

The exact number of anchors and width and height ratios that are used is determined by running k -means clustering on the dimensions of ground truth bounding boxes [5]. $k = 5$ was chosen in the original paper because it gives the best trade off between recall and complexity. In our implementation we use the 5 anchor boxes that give the lowest IOU on the VOC dataset.

2.3. Yolo loss function

The Yolo network uses a rather complex loss function. This function is comprised of multiple parts, designed to describe both the loss of the predicted bounding box(s) and class(es). The beauty of this loss function comes from the fact that only a single forward and backward pass of the network needs to be made in order for detection model to




Figure 3. Yolo Anchor Boxes [2].

improve. This is in contrast to several other detection models which require multiple passes of the image.

$$\begin{aligned} & \lambda_{\text{coord}} \sum_{i=0}^{S^2} \sum_{j=0}^B \mathbb{I}_{ij}^{\text{obj}} \left[(x_i - \hat{x}_i)^2 + (y_i - \hat{y}_i)^2 \right] \\ & + \lambda_{\text{coord}} \sum_{i=0}^{S^2} \sum_{j=0}^B \mathbb{I}_{ij}^{\text{obj}} \left[(\sqrt{w_i} - \sqrt{\hat{w}_i})^2 + (\sqrt{h_i} - \sqrt{\hat{h}_i})^2 \right] \\ & + \sum_{i=0}^{S^2} \sum_{j=0}^B \mathbb{I}_{ij}^{\text{obj}} (C_i - \hat{C}_i)^2 \\ & + \lambda_{\text{noobj}} + \sum_{i=0}^{S^2} \sum_{j=0}^B \mathbb{I}_{ij}^{\text{noobj}} (C_i - \hat{C}_i)^2 \\ & + \sum_{i=0}^{S^2} \mathbb{I}_{ij}^{\text{obj}} \sum_{c \in \text{classes}} (p_i(c) - \hat{p}_i(c))^2 \end{aligned}$$

\mathbb{I}^{obj} indicates a mask value used to ignore objects that are not present within a given cell. i.e. if an object appears \mathbb{I}^{obj} produces a 1. On the other hand if no object is present it produces a 0. In the above loss, $\sum_{i=0}^{S^2}$ is a sum over the 13×13 grid cells and $\sum_{j=0}^B$ is a sum over all 5 anchors.

Bounding Box Coordinate Loss

The two top most equations of the loss function describe the, amount of inaccuracy between the predicted bounding boxes x/y offsets and width/height ratios, when compared against that of the true bounding box.

Classification Loss

The next two equations in the loss penalise incorrect classification and superfluous classification respectively. This part pushes the network to the correct class for each object.

Confidence Score Loss

The final equation simply describes the difference between the predicted "confidence score" and the true class probability.

Lambda Hyper-parameters

There are two lambda values, coord and noobj. These values are used adjust the relative weight of the loss. increasing λ_{coord} will increase the bounding box's contribution to the overall model loss. Where as λ_{noobj} controls the contribution of "false objects" to the loss. As with many deep learning models, Yolo's loss is the most complex and most important part of the model.

2.4. Training

One of the reasons for the success of the Yolo network is the fact that an existing network trained for classification can be modified to perform object detection. In the original implementation the authors trained a network on ImageNet for about a week before modifying it to perform object detection. Pretraining for classification allows the network to start with reasonable weights before attempting the more complicated task of bounding box predictions. A classification network is modified for detection by removing the last convolutional layer and adding more convolutional layers with a shape dependent on the number of object classes. Another innovative aspect of Yolo9000 is that they can train for detection or classification, and backpropagate differently depending on what the input image is. For example, an input image without any bounding boxes will only update with respect to the classification part of the loss.

2.5. Results

Yolo9000 is able to perform "fast, accurate detection". Figure 4 shows the original implementations mAP and frames per second on the VOC2007 dataset in comparison to several other state of the art object detection systems. The speed of Yolo9000 contributes a lot to its success. There is even progress being made to implement the Yolo network on embedded systems. Using a similar network configuration with fewer parameters, Fast Yolo was able to achieve 18FPS on an embedded system [6].

3. Contributions of this project

For this paper we implemented a copy of the Yolo9000 network. Said copy contains a partial implementation of the




Figure 4. Yolo9000 Mean Average Precision and Frames per Second on VOC2007 [5]

Yolo9000 network, loss function, and optimiser. This implementation adds to the body of knowledge, pertaining to the implementation of Yolo in the TensorFlow deep learning framework. Unlike the original paper which trained not only on VOC 2012, COCO, and ImageNet [5], we only trained our network on the VOC 2012 data set. This decision was made not only to reduce the outrageous week long training times reported in the Yolo paper [3] but, also to reduce the chance of a mistake in the design of the network. As for the rest of the network architecture, our model was designed to be as close as possible to the original model put forward in the Yolo9000 paper. However, this duplication of the model did not go as planned, yielding mostly incorrect predictions due to not enough training time. Even with our subpar results, we have produced a reference implementation that will be helpful to future researchers and students pursuing the goal of running a Yolo9000 network using the TensorFlow framework. Perhaps, with a little luck, they may succeed where so many others have failed.

3.1. Training

The computational resources available to us did not allow us to pretrain a classification convolutional network on ImageNet. Instead our network started from scratch. We trained using the VOC2012 data which consists of 5700 training images and objects belonging to 20 classes. We used a Nvidia Tesla with 12GB of VRAM as that was required for a batch size of 32 using batch normalization. Our preprocessing of the images was fairly simple and consisted of resizing the images to 416x416 and performing random horizontal flipping of objects inside the bounding boxes. Finally, we used a batch size of 32, 45000 iterations, and a learning rate of 3e-5. Adam optimization was also used so the learning rate would slowly decrease with momentum. However, we ran into several issues during training which

required us to restart the training process. We expected the training to take 48 hours before any reasonable results appeared, but with the restarts we were unable to accomplish this with a bug free network.

3.2. Results

Our preliminary results can be seen in Figure 5. This figure shows the first bounding box predictions our network made. This was accompanied by a large spike in loss, as seen in Figure 6. We were initially concerned with this spike, but upon closer inspection we realised it corresponded to the network finally making bounding box predictions. The loss spiked because the IOU of these predictions was finally being taken into account.




Figure 5. First bounding box predictions of our network during training. [Left] ground truth [Right] predictions




Figure 6. Loss of our network after 28k iterations

3.3. Issues

The most challenging part of implementing Yolo9000 is getting the training data into the correct format and implementing the loss function. The use of anchor boxes is ingenious, yet difficult to implement. The loss function expects the input data to be in a specific format in a specific shape. The bounding boxes imported from the VOC2012 dataset is in the $[xmin, ymin, xmax, ymax]$ format, but needs to be converted to offsets relative to its closest grid cell and anchor box with the lowest IOU. The long training times made catching bugs with our implementation difficult.

One of the main issues we had was training time. Since training the network to make somewhat accurate predictions

takes upwards of 48 hours, any modifications we made, big or small, would take hours to be noticeable. On a few occasions we found small bugs in our implementation when we were already multiple hours into training. The only option was to stop training and start again. Another issue we had was the computational resources required to train the network. Training the network with a single image in a batch was almost impossible and painfully slow on anything but a powerful GPU. This increased the time for the compile-run-test loop and prevented us from quickly testing and iterating on our code.

Another issue we ran into was the size of the Tensorflow summaries. While developing we thought it was a good idea to have the preprocessed images along with their bounding boxes appear in Tensorboard. Initially this let us know we were preprocessing the images correctly. However, after 35k iterations of training the Tensorflow summary file grew to an unmanageable size (20GB) and forced us to stop training. These issues were learned from and will be avoided in the future.

4. Discussion and future directions

Our implementation of Yolo9000 is currently erroneous and does not produce meaningful results, as can be seen in figure 5. That being said, with a full length training period, we believe that we will begin to see performance comparable to that reported by Redmon et al. With a full training period we should end up with a state of the art object detection system. However, that is not to say that the Yolo system is without weakness. The main benefit of Yolo is the speed at which it can evaluate an image. This speed, however, is favoured at the cost of accuracy. Other versions of Yolo, such as Tiny Yolo, are designed to be even faster with less accuracy. The various forms of Yolo are thus well-suited to applications which require rapid evaluation, such as when processing video streams. However, this solution is likely not the ideal option if your intended application favours accuracy over speed. Another potential deterrent from using the Yolo system is the relatively long required pre-training. To achieve the reported performance, Redmon et al. first pre-trained their network for classification on ImageNet for around a week. From this, a modified version of the pretrained network is trained for detection, with the input resolution being doubled[4]. This results in a very time consuming training process, particularly when compared to networks which do not require pretraining.

Overall, the Yolo object detection system provides accurate, general purpose object classification and detection at an extremely fast speed. Though not perfect for all situations, this system has the potential to be used in a plethora of applications.

References

- [1] Tensorflow. <https://www.tensorflow.org/>. (Accessed on 04/06/2018).
- [2] C. Bourez. Bounding box object detectors: understanding yolo, you look only once, 2017.
- [3] J. Redmon, S. K. Divvala, R. B. Girshick, and A. Farhadi. You only look once: Unified, real-time object detection. *CoRR*, abs/1506.02640, 2015.
- [4] J. Redmon, S. K. Divvala, R. B. Girshick, and A. Farhadi. You only look once: Unified, real-time object detection. *CoRR*, abs/1506.02640, 2015.
- [5] J. Redmon and A. Farhadi. YOLO9000: better, faster, stronger. *CoRR*, abs/1612.08242, 2016.
- [6] M. J. Shafiee, B. Chywl, F. Li, and A. Wong. Fast YOLO: A fast you only look once system for real-time embedded object detection in video. *CoRR*, abs/1709.05943, 2017.
- [7] F. Shaikh. 10 advanced deep learning architectures data scientists should know!, 2018.